

Multi-level Model

Bias in PERSUADE

```
library(sjPlot)
library(lme4)
library(car) # level names in contrasts
library(stargazer) # LaTeX tables
library(emmeans)
library(performance) # ICC
library(tidyverse)
library(forcats)

dfp <- read.csv("../data/persuade-predictions-long.csv",
                na.strings = c("", " ", "NA")) %>%
  select(-c(full_text, X, assignment, source_text)) %>%
  mutate_if(is.character, as.factor) %>%

  # drop American Indian/Alaskan Native group because there are < 200 samples
  filter(race_ethnicity != "American Indian/Alaskan Native") %>%

  # remove unused levels (American Indian/Alaskan Native)
  droplevels() %>%

  # simplify level and variable names
  mutate(
    Transformer=fct_recode(model,
                          Black="baa_pred",
                          Control="control_pred",
                          Hispanic="hl_pred",
                          White="w_pred"),
    race=fct_recode(race_ethnicity,
                   asian="Asian/Pacific Islander",
                   black="Black/African American",
```

```

        hisp="Hispanic/Latino",
        other="Two or more races/Other",
        white="White")
    ) %>%

# calculate root mean squared error
mutate(rmse=sqrt(squared_error)) %>%

# calculate percent error
mutate(percent_error = (pred_error/holistic_score_adjudicated)*100) %>%

# collapse NaNs into negative level of binary factors
mutate(
  is_ell=fct_collapse(addNA(ell), Yes = "Yes", No = c("No", NA)),
  is_disadvantaged=fct_collapse(addNA(economically_disadvantaged),
                                Yes = "Economically disadvantaged",
                                No = c("Not economically disadvantaged",
                                       NA)
                                ),
  has_disability=fct_collapse(addNA(student_disability_status),
                              Yes = "Identified as having disability",
                              No = c("Not identified as having disability",
                                     NA)
                              )
  ) %>%

# set the first level, which will be used as "reference"
mutate(
  is_disadvantaged=fct_relevel(is_disadvantaged, "No"),
  has_disability=fct_relevel(has_disability, "No"),
  Transformer=fct_relevel(Transformer, "Control")
) %>%

mutate(
  Direction = case_when(
    pred_error >= 0 ~ "Positive",
    pred_error < 0 ~ "Negative"
  ),
  Direction = factor(Direction, levels=c("Positive","Negative"))
)

```

```
# configure contrasts for race and source
options(decorate.contr.Sum = c("", ""))
contrasts(dfp$race) = contr.Sum(levels(dfp$race))
contrasts(dfp$source) = contr.Sum(levels(dfp$source))

summary(dfp)
```

essay_id		holistic_score_1	holistic_score_2
2021000047:	4	Min. :1.000	Min. :1.000
2021000136:	4	1st Qu.:2.000	1st Qu.:2.000
2021000241:	4	Median :3.000	Median :3.000
2021000462:	4	Mean :3.364	Mean :3.388
2021000730:	4	3rd Qu.:4.000	3rd Qu.:4.000
2021000764:	4	Max. :6.000	Max. :6.000

(Other) :19156

holistic_score_adjudicated	source
Min. :1.000	Florida :2436
1st Qu.:3.000	Georgia Virtual: 876
Median :3.000	Indiana :6384
Mean :3.405	NCES :3120
3rd Qu.:4.000	Virginia :6364
Max. :6.000	

	prompt_name	task	gender
Distance learning	:1872	Independent :10360	F:9612
Summer projects	:1640	Text dependent: 8820	M:9568
Facial action coding system	:1576		
Mandatory extracurricular activities:	1532		
Exploring Venus	:1448		
Driverless cars	:1384		
(Other)	:9728		

grade	ell	race_ethnicity
Min. : 6.000	No :16408	Asian/Pacific Islander :3836
1st Qu.: 8.000	Yes : 1812	Black/African American :3836
Median :10.000	NA's: 960	Hispanic/Latino :3836
Mean : 9.244		Two or more races/Other:3836
3rd Qu.:10.000		White :3836
Max. :12.000		

economically_disadvantaged

Economically disadvantaged :7660
 Not economically disadvantaged:8120
 NA's :3400

	student_disability_status	model
Identified as having disability	: 2300	baa_pred :4795
Not identified as having disability	:16004	control_pred:4795
NA's	: 876	hl_pred :4795
		w_pred :4795

	prediction	pred_error	squared_error	Transformer
Min.	:0.795	Min. :-2.2323	Min. : 0.00000	Control :4795
1st Qu.:	2.709	1st Qu.: -0.2752	1st Qu.: 0.03495	Black :4795
Median	:3.453	Median : 0.1179	Median : 0.14918	Hispanic:4795
Mean	:3.527	Mean : 0.1212	Mean : 0.36447	White :4795
3rd Qu.:	4.285	3rd Qu.: 0.4887	3rd Qu.: 0.42920	
Max.	:5.884	Max. : 3.4119	Max. :11.64123	

	race	rmse	percent_error	is_ell	is_disadvantaged
asian:3836	Min.	:0.00003	Min. :-49.861	No :17368	No :11520
black:3836	1st Qu.:	0.18696	1st Qu.: -6.937	Yes: 1812	Yes: 7660
hisp :3836	Median	:0.38623	Median : 3.677		
other:3836	Mean	:0.47062	Mean : 9.335		
white:3836	3rd Qu.:	0.65514	3rd Qu.: 16.232		
	Max.	:3.41193	Max. :341.192		

has_disability	Direction
No :16880	Positive:11184
Yes: 2300	Negative: 7996

Summary statistics

```
dfp %>%
  filter(Transformer == "Control") %>%
  summarize(
    Score = mean(holistic_score_adjudicated),
    N = n(),
    ELLs = sum(as.numeric(is_ell)),
    Disability = sum(as.numeric(has_disability)),
    Disadvantaged = sum(as.numeric(is_disadvantaged)),
    .by=c(race_ethnicity)
  ) %>%
  arrange(race_ethnicity)
```

	race_ethnicity	Score	N	ELLs	Disability	Disadvantaged
1	Asian/Pacific Islander	3.910323	959	1117	1018	1229
2	Black/African American	3.103233	959	997	1086	1464
3	Hispanic/Latino	3.104275	959	1193	1068	1524
4	Two or more races/Other	3.453597	959	970	1104	1344
5	White	3.455683	959	971	1094	1149

Model from Research Question 1

Just adding in transformer model as a predictor and a random intercept for essay ID. This model has poor fit, with many useless parameters. Prompt_name is no longer useful as a grouping variable, so we can drop it.

```
mod.rq1 <- lmer(pred_error
  ~ trf
  + gender
  + race*is_ell
  + is_disadvantaged*has_disability
  + is_disadvantaged*is_ell
  + has_disability*is_ell
  + (1|prompt_name)
  + (1|essay_id),
  data=dfp
)
tab_model(mod.rq1, show.aic=TRUE, show.reflavl=TRUE, prefix.labels="varname", p.adjust = "B")
```

Save Table

```
coefs <- summary(fit)$coefficients
cat(rownames(coefs), sep="\\", "\\")

labels <- c("(Intercept)",
            "Transformer[Black]",
            "Transformer[Hispanic]",
            "Transformer[White]",
            "Black/African American",
            "Hispanic/Latino",
            "Two or more races/Other",
            "White",
            "English language learner",
            "Identified as having disability",
            "Economically disadvantaged",
            "Transformer[Black]:Black/African American",
            "Transformer[Hispanic]:Black/African American",
            "Transformer[White]:Black/African American",
            "Transformer[Black]:Hispanic/Latino",
            "Transformer[Hispanic]:Hispanic/Latino",
            "Transformer[White]:Hispanic/Latino",
            "Transformer[Black]:Two or more races/Other",
            "Transformer[Hispanic]:Two or more races/Other",
            "Transformer[White]:Two or more races/Other",
            "Transformer[Black]:White",
            "Transformer[Hispanic]:White",
            "Transformer[White]:White",
            "Transformer[Black]:English language learner",
            "Transformer[Hispanic]:English language learner",
            "Transformer[White]:English language learner",
            "Transformer[Black]:Identified as having disability",
            "Transformer[Hispanic]:Identified as having disability",
            "Transformer[White]:Identified as having disability",
            "Transformer[Black]:Economically disadvantaged",
            "Transformer[Hispanic]:Economically disadvantaged",
            "Transformer[White]:Economically disadvantaged")

tab_model(fit,
          title = "Prediction Error Regressed on Demographic Variables",
          dv.labels = "Prediction Error",
```

```

show.aic=TRUE,
show.reflvl=TRUE,
prefix.labels="varname",
p.adjust = "BH",
pred.labels=labels,
file = "../results/RQ2.html")

```

Tear-down

We would be left with very little in the model.

```

fit.min <- lmer(pred_error
  ~ trf
  # + race # increases AIC
  # + source # increases AIC
  # + is_disadvantaged # increases AIC
  + is_ell
  + has_disability
  # + gender # increases AIC
  # + race*is_disadvantaged # increases AIC
  # + race*is_ell # increases AIC
  # + race*has_disability # increases AIC
  # + race*source # increases AIC
  # + race*gender # increases AIC
  # + race*grade # increases AIC
  # + is_disadvantaged*has_disability # increases AIC
  # + is_disadvantaged*is_ell # increases AIC
  # + has_disability*is_ell # increases AIC
  # + gender*has_disability # increases AIC
  # + gender*is_disadvantaged # increases AIC
  # + gender*is_ell # small decrease (-1) to AIC
  + (1 | essay_id),
  data=dfp
)

```

```

tab_model(fit.min, show.aic=TRUE, show.reflvl=TRUE, prefix.labels="varname", p.adjust = "B

```

Build-up: Interactions

While race did not help model fit on its own, its interaction with transformer model substantially improves model fit. The same pattern holds for interactions between transformer model and other demographic variables.

```
fit <- lmer(pred_error
  ~ Transformer
  + Transformer*race_ethnicity
  + Transformer*is_ell
  + Transformer*has_disability
  + Transformer*is_disadvantaged
  + (1 | essay_id),
  REML = TRUE,
  data=dfp
)

# tab_model(fit, show.aic=TRUE, show.reflvl=TRUE, prefix.labels="varname", p.adjust = "BH")
```

Deviance Tests

These require non-restricted ML estimations.

```
fit.ml <- lmer(pred_error
  ~ Transformer
  + Transformer*race
  # + Transformer*is_ell
  + Transformer*has_disability
  + Transformer*is_disadvantaged
  + (1 | essay_id),
  REML = FALSE,
  data=dfp
)

fit.ell <- lmer(pred_error
  ~ Transformer
  + Transformer*race
  + Transformer*is_ell
  + Transformer*has_disability
  + Transformer*is_disadvantaged
```



```

+ (1 | essay_id),
REML = FALSE,
data=dfp
)

1-pchisq(deviance(fit.ml) - deviance(fit.ell), df=2)

```

```
[1] 0
```

RMSE

RMSE is more interpretable. Possibly a separate research question?

```

fit.rmse <- lmer(rmse
  ~ Transformer
  + Transformer*race_ethnicity
  + Transformer*is_ell
  # + Transformer*has_disability
  + Transformer*is_disadvantaged
  + (1 | essay_id),
  REML = TRUE,
  data=dfp
)

tab_model(fit.rmse, show.aic=TRUE, show.reflvl=TRUE, prefix.labels="varname", p.adjust = "

```

Interaction Plots

```

# c("Asian/Pacific Islander",
#   "Black/African American",
#   "Hispanic/Latino",
#   "Two or more races/Other",
#   "White")

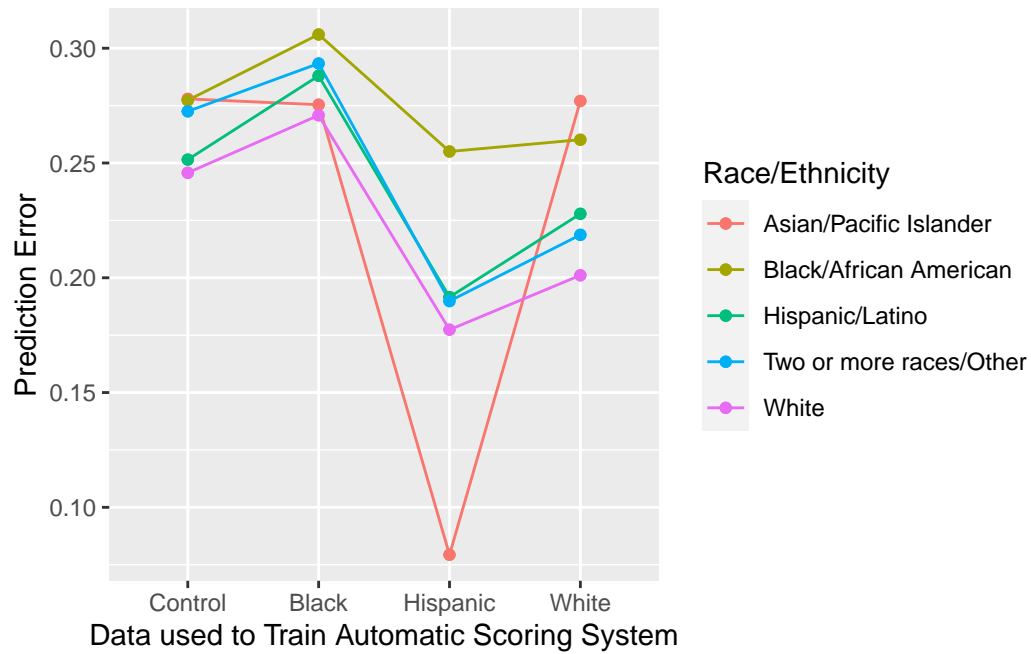
emmip(fit, race_ethnicity ~ Transformer,
      mode="asympt",
      type = "scale",
      style = "factor",

```

```

tlab = "Race/Ethnicity",
ylab = "Prediction Error",
xlab = "Data used to Train Automatic Scoring System"
)

```



```

# ggsave("../results/transformer_race.png")

```